Spring 2021: Numerical Analysis
Assignment 2 (due Monday March 8th 2pm)

1. **[Compound interest, 8pts]** For a yearly interest rate $0 < r < 1$ compounded over n intervals, an amount of money C grows to be

$$f(C, r, n) = C \left(1 + \frac{r}{n}\right)^n \tag{1}$$

after one year. Let $C = 1$ and $r = 0.025$. If n is large, there may be loss of digits when evaluating this using finite-precision arithmetic.

(a) [2pts] If n is *extremely* large, say $n = 10^{16}$, in IEEE double precision arithmetic (try it in Matlab),

$$f\left(1.0, 0.025, 10^{16}\right) = 1.0, \tag{2}$$

when in fact,

$$\begin{aligned} f(1.0, 0.025, 10^{16}) &\approx \lim_{n \to \infty} \left(1 + \frac{0.025}{n}\right)^n \\ &= e^{0.025} \\ &\approx 1.025315\ldots \end{aligned} \tag{3}$$

What happened?

(b) [2pts] To compute $f(C, r, n)$ without roundoff problems in Matlab, compute first $\ln f$ using the (magic!) built-in Matlab function $log1p$ which computes $\ln(1+x)$ without loosing digits even for very small x, and then compute f from its logarithm. Write down the formulas used. Try this for $n = 10^{16}$. From now on take $n = 10^8$.

(c) [2pts] Using the result from part (b), how many digits of accuracy do you get for f with direct evaluation of (1)?

(d) [1pts] For large n, we can just use the approximation $f(C, r, n) = Ce^r$. How many digits of accuracy do you get with this approximation?

(e) [1pts] An improved approach for large n is to compute a few terms in the Taylor series expansion (not a trivial calculation per se),

$$(1 + rx)^{1/x} = e^r \left[1 - \frac{r^2}{2}x + O\left(x^2\right)\right],$$

and then use this approximation for small x. How many digits of accuracy do you get using this approach?

Don't just report answers, explain how you computed this.

2. **[Backward substitution implementation, 5pts]** [3pts] Write a code for backward substitution to solve systems of the form $Ux = b$, i.e., write a function `x = backward(A,b)`, which expects as inputs an upper triangular matrix $U \in \mathbb{R}^{n \times n}$, and a right hand side vector $b \in \mathbb{R}^n$, which returns the solution vector $x \in \mathbb{R}^n$. The function should find the size n from the vector b and also check if the matrix and the vector sizes are compatible before it starts to solve the system. Apply your program for the computation of for $x \in \mathbb{R}^4$, with

$$U = \begin{bmatrix} 1 & 2 & 6 & -1 \\ 0 & 3 & 1 & 0 \\ 0 & 0 & 4 & -1 \\ 0 & 0 & 0 & 2 \end{bmatrix}, \quad b = \begin{bmatrix} -1 \\ -3 \\ -2 \\ 4 \end{bmatrix}.$$

   [2pts] How do you know that your code is working correctly?

3. **[LU factorization of tridiagonal matrix, 6pt]** Given is a tridiagonal matrix, i.e., a matrix with nonzero entries only in the diagonal, and the first upper and lower subdiagonals:

$$A = \begin{bmatrix} a_1 & c_1 & & & \\ b_1 & a_2 & c_2 & & \\ & \ddots & \ddots & \ddots & \\ & & b_{n-2} & a_{n-1} & c_{n-1} \\ & & & b_{n-1} & a_n \end{bmatrix}.$$

   Assuming that A has an LU decomposition $A = LU$ with

$$L = \begin{bmatrix} 1 & & & \\ d_1 & 1 & & \\ & \ddots & \ddots & \\ & & d_{n-1} & 1 \end{bmatrix}, \quad U = \begin{bmatrix} e_1 & f_1 & & \\ & \ddots & \ddots & \\ & & e_{n-1} & f_{n-1} \\ & & & e_n \end{bmatrix},$$

   derive iterative expressions for d_i, e_i and f_i (i.e., how to compute the value for $i+1$ from the values at i, and how to start for $i = 1$).

   Hint: You could check your answer by implementing the formulas in code and checking that $LU = A$ in Matlab for some specific example.

4. **[Inverse matrix computation, 8pts]** Let us use the LU-decomposition to compute the inverse of a matrix[1].

   (a) [2pts] Describe an algorithm that uses the LU-decomposition of an $n \times n$ matrix A for computing A^{-1} by solving n systems of equations (one for each unit vector).

---

[1]This also illustrates that computing a matrix inverse is significantly more expensive than solving a linear system. That is why to solve a linear system, you should *never* use the inverse matrix!

(b) [2pts] Calculate the floating point operation count of this algorithm. It is OK to use estimates from class/worksheets but write them down so the grader knows what you are doing.

(c) [4pts] Improve the algorithm by taking advantage of the structure (i.e., the many zero entries) of the right-hand side. What is the new algorithm's floating point operation count?

[*Hint:* Consider splitting the solution vector for the k-th equation from part (a) into two pieces, and solve for each piece separately, on paper or using forward/back substitution.]

5. **[Stability of the Gaussian elimination, 8pts]**

Consider the linear system
$$\boldsymbol{A}\boldsymbol{x} = \boldsymbol{b}, \tag{4}$$

where $\boldsymbol{A}$ is an $n \times n$ matrix that has ones on the diagonal, minus ones below the diagonal, and ones in the last column, with all other entries zero. For example, when $n = 5$, we have
$$\boldsymbol{A} = \begin{pmatrix} 1 & 0 & 0 & 0 & 1 \\ -1 & 1 & 0 & 0 & 1 \\ -1 & -1 & 1 & 0 & 1 \\ -1 & -1 & -1 & 1 & 1 \\ -1 & -1 & -1 & -1 & 1 \end{pmatrix}.$$

(a) [3pts] Prove that A is invertible for any n, by induction. [Hint: Perform a *column operation* on $\boldsymbol{A}$ to eliminate the reduce it to a smaller matrix of size $n - 1$ and ask whether that smaller matrix is invertible under the induction hypothesis.]

(b) [3pts] Now consider the matrix $\boldsymbol{A}$ for some unspecified (arbitrary) n. Perform Gaussian elimination on $\boldsymbol{A}$ to obtain the upper triangular matrix $\boldsymbol{U}$ appearing in the LU factorization $\boldsymbol{A} = \boldsymbol{LU}$. What is $\max_{i,j} |u_{i,j}|$ as a function of n?

(c) [2pts] For large n, e.g., $n = 2000$, what problems can you envision if you try to solve (4) using Gaussian elimination on a computer? Explain.
[*Note: This is one rare example matrix for even Matlab will fail to solve a linear system correctly even though the matrix is well-conditioned, see discussion in Section 7.5 of Practice textbook.*]

6. **[Matrix square root, 6pts]** Newton's method for finding roots can be extended to *matrix-valued functions* as well. Here you will devise a Newton method (i.e., generalize the Babylonian method) to compute the square root of a matrix. If it exists, the square root of a real symmetric $n \times n$ matrix $\boldsymbol{A}$ is another *real square symmetric* matrix $\boldsymbol{X}$ such that
$$\boldsymbol{X}^T \boldsymbol{X} = \boldsymbol{A} \tag{5}$$

Just like the square root of even a positive number is not unique, the matrix square is not unique (one can roughly think of having to choose n signs, as we will revisit in a future homework once we cover eigenvalue decompositions).

(a) [2pts] In class/worksheet we used derivatives to obtain Newton's method. Instead of computing derivatives of matrix-valued functions, however, it is useful to think of computing derivatives from a *linearization* of the function around a given value (this allows to generalize the notion of a derivative and makes it easier to compute in some cases). Set $X = X + \delta X$ in (5) and keep only the terms that are linear in the 'perturbation" δX. Use this to write down an equation for δX.

(b) [2pts] The equation you obtained in part (a) can be solved explicitly for any n — can you explain why? It is OK if you assume a unique solution exists. Take $n = 2$ and write down the solution explicitly. [*Hint: It is always a good idea to check by plugging in specific numbers.*]

(c) [2pts] It would be nice to write down an explicit formula for the solution of the equation you got from part (a) for any n. Do this by assuming that the matrices X and δX commute, i.e., that

$$X\,(\delta X) = (\delta X)\,X. \tag{6}$$

[Hint: Recall that X is symmetric.].
Note: One can prove (6) holds at all iterations if the initial guess X_0 commutes with A; if interested, look at the paper "Newton's Method for the Matrix Square Root" by Nicholas Higham, freely available on the web.