

# Eigenvalues and Singular Values

**Aleksandar Donev**

*Courant Institute, NYU<sup>1</sup>*

*donev@courant.nyu.edu*

<sup>1</sup>Course MATH-UA.0252/MA-UY\_4424, Spring 2021

Spring 2021

# Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Uses of the SVD
- 4  $QR$  method for computing eigenvalues

# Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Uses of the SVD
- 4  $QR$  method for computing eigenvalues

# The need for iterative algorithms

- The eigenvalues are roots of the **characteristic polynomial** of **A**, which is generally of order  $n$ .
- According to Abel's theorem, there is no closed-form (rational) solution for  $n \geq 5$ .

**All eigenvalue algorithms must be iterative!**

- There is an important distinction between iterative methods to:
  - Compute **all eigenvalues** (similarity transformations). These are based on dense-matrix factorizations such as the  $QR$  factorization, with total cost  $O(n^3)$ .
  - Compute **only one or a few eigenvalues**, typically the smallest or the largest one (e.g., power method). These are similar to iterative methods for solving linear systems.

# Sparse Matrices

- Recall that for a diagonalizable matrix

$$\mathbf{A}^n = \mathbf{X}\mathbf{\Lambda}^n\mathbf{X}^{-1}$$

and **assume well-separated eigenvalues**  $|\lambda_1| > |\lambda_2| \geq |\lambda_3| \cdots |\lambda_n|$ , and that the columns of  $\mathbf{X}$  are normalized,  $\|\mathbf{x}_j\| = 1$ .

- For sparse matrices we sometimes only need to know a **few of the eigenvalues**/vectors, not all of them.
- Notably, knowing the eigenvector corresponding to the **smallest and largest (in magnitude) eigenvalues** is often most important (see Google Page Rank algorithm).

# Iterative Method

- Any **initial guess** vector  $\mathbf{q}_0$  can be represented in the linear basis formed by the eigenvectors

$$\mathbf{q}_0 = \mathbf{X}\mathbf{a}$$

- Recall iterative methods for linear systems: **Multiply a vector with the matrix  $\mathbf{A}$**  many times:

$$\mathbf{q}_{k+1} = \mathbf{A}\mathbf{q}_k$$

$$\mathbf{q}_n = \mathbf{A}^n \mathbf{q}_0 = (\mathbf{X}\mathbf{\Lambda}^n \mathbf{X}^{-1}) \mathbf{X}\mathbf{a} = \mathbf{X}(\mathbf{\Lambda}^n \mathbf{a})$$

# Power Method

- As  $n \rightarrow \infty$ , the **eigenvalue of largest modulus**  $\lambda_0$  will dominate,

$$\mathbf{\Lambda}^n = \lambda_1^n \text{Diag} \left\{ 1, \left( \frac{\lambda_2}{\lambda_1} \right)^n, \dots \right\} \rightarrow \text{Diag} \{ \lambda_1^n, 0, \dots, 0 \}$$

$$\mathbf{q}_n = \mathbf{X} (\mathbf{\Lambda}^n \mathbf{a}) \rightarrow \lambda_1^n \mathbf{X} \begin{bmatrix} a_1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \lambda_1^n \mathbf{x}_1$$

- Therefore the **normalized iterates** converge to the eigenvector:

$$\tilde{\mathbf{q}}_n = \frac{\mathbf{q}_n}{\|\mathbf{q}_n\|} \rightarrow \mathbf{x}_1$$

- The **Rayleigh quotient** converges to the eigenvalue:

$$r_A(\mathbf{q}_n) = \frac{\mathbf{q}_n^* \mathbf{A} \mathbf{q}_n}{\mathbf{q}_n \cdot \mathbf{q}_n} = \tilde{\mathbf{q}}_n^* \mathbf{A} \tilde{\mathbf{q}}_n \rightarrow \lambda_1$$

## An alternative derivation

$$\begin{aligned}
 q_0 &= \sum_n a_i \cdot x_i \\
 q_n &= A^n q_0 = \sum_n a_i (A^n x_i) \\
 &= \sum_n a_i (\lambda_i^n x_i) = \\
 &= \sum_n a_i \lambda_1^n \left( \frac{\lambda_i}{\lambda_1} \right)^n x_i \xrightarrow[n \rightarrow \infty]{\text{as}} a_1 \lambda_1^n x_1 + O\left(\left| \frac{\lambda_2}{\lambda_1} \right|^n\right)
 \end{aligned}$$

$$\Rightarrow \tilde{q}_n \rightarrow x_1 + O(\epsilon) x_2 + O(\epsilon) x_3 + \dots$$

where  $\epsilon = \left| \frac{\lambda_2}{\lambda_1} \right|^n$

Now  $\tilde{q}_n^* A \tilde{q}_n = x_1^* A x_1 + O(\epsilon^2) \sum_{i,j \neq 1} x_i^* A x_j$

$$\Rightarrow \hat{\lambda}_n = \lambda_1 + O(\epsilon^2)$$



# Power Iteration

Start with an initial guess  $\mathbf{q}_0$ , and then iterate:

- 1 Compute **matrix-vector product** and normalize it:

$$\mathbf{q}_k = \frac{\mathbf{A}\mathbf{q}_{k-1}}{\|\mathbf{A}\mathbf{q}_{k-1}\|}$$

- 2 Use Raleigh quotient to obtain **eigenvalue estimate**:

$$\hat{\lambda}_k = \mathbf{q}_k^* \mathbf{A} \mathbf{q}_k$$

- 3 **Test for convergence**: Evaluate the residual

$$\mathbf{r}_k = \mathbf{A}\mathbf{q}_k - \hat{\lambda}_k \mathbf{q}_k$$

and terminate if the residual norm is smaller than some tolerance, e.g., for tolerance  $\epsilon \ll 1$ ,

$$\|\mathbf{r}_k\| \approx \left| \lambda_1 - \hat{\lambda}_k \right| \leq \epsilon \hat{\lambda}_k.$$

# Convergence Estimates

- The normalized iterates converge to the eigenvector **linearly**:

$$\|\mathbf{q}_k - (\pm \mathbf{x}_1)\| = O\left(\left|\frac{\lambda_2}{\lambda_1}\right|^k\right)$$

- Typically the eigenvalue estimate converges **quadratically**:

$$\|\hat{\lambda}_k - \lambda_1\| \sim O\left(\left|\frac{\lambda_2}{\lambda_1}\right|^{2k}\right)$$

- The power method is fast when the **dominant eigenvalue is well-separated from the rest** (even if it is degenerate).
- This conclusion is rather general for all iterative methods:  
Convergence is good for **well-separated** eigenvalues, bad otherwise.

# Eigenvalues in MATLAB

- The **Schur decomposition** is provided by  $[U, T] = \text{schur}(A)$ .
- In MATLAB, sophisticated variants of the **QR algorithm** (LAPACK library) are implemented in the function *eig*:

$$\Lambda = \text{eig}(A)$$

$$[X, \Lambda] = \text{eig}(A)$$

- For large or sparse matrices, iterative methods based on the **Arnoldi iteration** (ARPACK library), can be used to obtain a **few** of the largest eigenvalues:

$$\Lambda = \text{eigs}(A, n_{\text{eigs}})$$

$$[X, \Lambda] = \text{eigs}(A, n_{\text{eigs}})$$

# Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Uses of the SVD
- 4  $QR$  method for computing eigenvalues

# Sensitivity (conditioning) of the SVD

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^*$$

- Since unitary matrices have unit 2-norm,

$$\|\delta\mathbf{\Sigma}\|_2 \approx \|\delta\mathbf{A}\|_2.$$

- The SVD computation is always **perfectly well-conditioned**!
- However, this refers to absolute errors: The **relative error** of small singular values will be large.
- The **power of the SVD** lies in the fact that it always exists and can be computed stably...but it is somewhat **expensive to compute**.

# Computing the SVD

- The SVD can be computed by performing an eigenvalue computation for the **normal matrix**  $\mathbf{A}^* \mathbf{A}$  (a positive-semidefinite matrix).
- This squares the condition number for small singular values and is **not numerically-stable**.
- Instead, modern algorithms use an algorithm based on computing eigenvalues / eigenvectors using the  $QR$  factorization.
- The cost of the calculation is  $\sim O(mn^2)$ , of the same order as eigenvalue calculation if  $m \sim n$ .

# Reduced SVD

## The **full (standard) SVD**

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^* = \sum_{i=1}^p \sigma_i \mathbf{u}_i \mathbf{v}_i^*$$

$$[m \times n] = [m \times m] [m \times n] [n \times n],$$

is in practice often computed in **reduced (economy) SVD** form, where  $\mathbf{\Sigma}$  is  $[p \times p]$ :

$$[m \times n] = [m \times n] [n \times n] [n \times n] \quad \text{for } m > n$$

$$[m \times n] = [m \times m] [m \times m] [m \times n] \quad \text{for } n > m$$

This contains all the information as the full SVD but can be **cheaper to compute** if  $m \gg n$  or  $m \ll n$ .

## In MATLAB

- $[U, \Sigma, V] = \text{svd}(A)$  for **full SVD**, computed using a QR-like method.
- $[U, \Sigma, V] = \text{svd}(A, 'econ')$  for **economy SVD**.
- The **least-squares solution** for square, overdetermined, underdetermined, or even rank-deficient systems can be computed using *svd* or *pinv* (pseudo-inverse, see homework).
- The  $q$  largest singular values and corresponding approximation can be computed efficiently for **sparse matrices** using

$$[U, \Sigma, V] = \text{svds}(A, q).$$



# Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Uses of the SVD
- 4  $QR$  method for computing eigenvalues

# Rank-Revealing Properties

- Assume the rank of the matrix is  $r$ , that is, the dimension of the range of  $\mathbf{A}$  is  $r$  and the dimension of the null-space of  $\mathbf{A}$  is  $n - r$  (recall the fundamental theorem of linear algebra).
- The SVD is a **rank-revealing** matrix factorization because only  $r$  of the singular values are nonzero,

$$\sigma_{r+1} = \cdots = \sigma_p = 0.$$

- The left singular vectors  $\{\mathbf{u}_1, \dots, \mathbf{u}_r\}$  form an **orthonormal basis for the range** (column space, or image) of  $\mathbf{A}$ .
- The right singular vectors  $\{\mathbf{v}_{r+1}, \dots, \mathbf{v}_n\}$  form an **orthonormal basis for the null-space** (kernel) of  $\mathbf{A}$ .

# The matrix pseudo-inverse

- For square non-singular systems,  $\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$ . Can we generalize the matrix inverse to non-square or rank-deficient matrices?
- Yes: **matrix pseudo-inverse** (Moore-Penrose inverse):

$$\mathbf{A}^\dagger = \mathbf{V}\mathbf{\Sigma}^\dagger\mathbf{U}^*,$$

where

$$\mathbf{\Sigma}^\dagger = \text{Diag} \{ \sigma_1^{-1}, \sigma_2^{-1}, \dots, \sigma_r^{-1}, 0, \dots, 0 \}.$$

- In numerical computations very small singular values should be considered to be zero (see homework).
- The least-squares solution to over- or under-determined linear systems  $\mathbf{Ax} = \mathbf{b}$  can be obtained from:

$$\mathbf{x} = \mathbf{A}^\dagger\mathbf{b}.$$

# Proof of Least-Squares (1)

$$\min_x \|Ax - b\|_2 \leftarrow \text{LEAST SQUARES}$$

SUCH THAT  $\|x\|_2$  IS MINIMAL

---

$$(Ax - b)^*(Ax - b) = x^*(A^*A)x - 2x^*A^*b + \dots$$


---

Using SVD:  $x^*A^*Ax = x^*V\Sigma^2V^*x$

AND  $x^*A^*b = b^*(Ax) = (U^*b)\underline{\underline{(V^*x)}}$

Denoting 
$$\begin{cases} V^*x = w \leftarrow \text{NEW VARIABLE} \\ U^*b = c \leftarrow \text{CONSTANT} \end{cases}$$

(1)

## Proof of Least-Squares (2)

$$\begin{aligned}
 \|Ax - b\|_2^2 &= w^* \Sigma^2 w - 2 c^* \Sigma w + \dots \\
 &= \sum_{i=1}^r \sigma_i^2 w_i^2 - 2 \sum_{i=1}^r (\sigma_i w_i) c_i^* \\
 &= \sum_{i=1}^r |\sigma_i w_i - c_i|^2 + \text{CONSTANTS}
 \end{aligned}$$

which is minimized if

$$\sigma_i w_i = c_i \Rightarrow w_i = \frac{c_i}{\sigma_i}$$

or

$$(V^* x)_i = \frac{(U^* b)_i}{\sigma_i}, \quad i \leq r$$

(2)

## Proof of Least-Squares (3)

How ABOUT  $w_{r+1}, \dots, w_m$  ?

$$\begin{aligned} \|X\|_2^2 &= \|Vw\|_2^2 = w^*(V^*V)w \\ &= \|w\|_2^2 = \sum |w_i|^2 \end{aligned}$$

SO THE NORM OF  $X$  IS MINIMIZED  
IF THE NORM OF  $w$  IS MINIMIZED.

$$\Rightarrow \boxed{w_{r+1} = \dots = 0}$$

$$\begin{aligned} \Rightarrow X &= Vw = V \Sigma^+ c = \\ &= (V \Sigma^+ u^*) b = A^+ b \end{aligned}$$

QED (3)

# Low-rank approximations

- The SVD is a decomposition into **rank-1 outer product matrices**:

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^* = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^* = \sum_{i=1}^r \mathbf{A}_i$$

- The rank-1 components  $\mathbf{A}_i$  are called **principal components**, the most important ones corresponding to the larger  $\sigma_i$ .
- Ignoring all singular values/vectors except the first  $q$ , we get a **low-rank approximation**:

$$\mathbf{A} \approx \hat{\mathbf{A}}_q = \mathbf{U}_q \mathbf{\Sigma}_q \mathbf{V}_q^* = \sum_{i=1}^q \sigma_i \mathbf{u}_i \mathbf{v}_i^*.$$

- Theorem: This is the **best approximation** of rank- $q$  in the Euclidian and Frobenius norm:

$$\left\| \mathbf{A} - \hat{\mathbf{A}}_q \right\|_2 = \sigma_{q+1}$$

# Applications of SVD

- **Statistical analysis** (e.g., DNA microarray analysis, clustering), often called **Principal Component Analysis (PCA)**
- Data **compression** (e.g., image compression, explained next).
- **Feature extraction**, e.g., face or character recognition (see Eigenfaces on Wikipedia).
- **Latent semantic indexing** for context-sensitive searching (see Wikipedia).
- **Noise reduction** (e.g., weather prediction).



# Image Compression

```
>> A=rgb2gray(imread('basket.jpg'));  
>> imshow(A);  
>> [U,S,V]=svd(double(A));  
>> r=25; % Rank-r approximation  
>> Acomp=U(:,1:r)*S(1:r,1:r)*(V(:,1:r))';  
>> imshow(uint8(Acomp));
```

# Compressing an image of a basket

We used only 25 out of the  $\sim 400$  singular values to construct a rank 25 approximation:



# Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Uses of the SVD
- 4 *QR* method for computing eigenvalues

# Estimating all eigenvalues / eigenvectors

- Iterative methods akin the power method are not suitable for estimating all eigenvalues.
- Basic idea: Build a sequence of matrices  $\mathbf{A}_k$  that all share eigenvalues with  $\mathbf{A}$  via **similarity transformations**:

$$\mathbf{A}_{k+1} = \mathbf{P}^{-1} \mathbf{A}_k \mathbf{P}, \text{ starting from } \mathbf{A}_1 = \mathbf{A}.$$

- The goal is to reduce the matrix  $\mathbf{A}_{k \rightarrow \infty}$  to as close to diagonal as possible.
- A numerically stable and good way to do this is to use the QR factorization:

$$\mathbf{A}_k = \mathbf{Q}_{k+1} \mathbf{R}_{k+1}$$

$$\mathbf{A}_{k+1} = \mathbf{Q}_{k+1}^{-1} \mathbf{A}_k \mathbf{Q}_{k+1} = (\mathbf{Q}_{k+1}^{-1} \mathbf{Q}_{k+1}) \mathbf{R}_{k+1} \mathbf{Q}_{k+1} = \mathbf{R}_{k+1} \mathbf{Q}_{k+1}.$$

# The basic QR method

- The behavior of the QR iteration can be understood most transparently as follows [following Trefethen and Bau]:
- Observation: The range of the matrix  $\mathbf{A}^k$  converges to the space spanned by the eigenvectors of  $\mathbf{A}$ , with the eigenvectors corresponding to the largest eigenvalues dominating as  $k \rightarrow \infty$  (so this is ill-conditioned).
- Recall: The columns of  $\mathbf{Q}$  in  $\mathbf{A} = \mathbf{QR}$  form an **orthonormal basis** for the range of  $\mathbf{A}$ .
- Idea: Form a **well-conditioned basis for the eigenspace** of  $\mathbf{A}$  by factorizing:

$$\mathbf{A}^k = \tilde{\mathbf{Q}}_k \tilde{\mathbf{R}}_k$$

and then calculate

$$\mathbf{A}_k = \tilde{\mathbf{Q}}_k^{-1} \mathbf{A} \tilde{\mathbf{Q}}_k = \tilde{\mathbf{Q}}_k^* \mathbf{A} \tilde{\mathbf{Q}}_k.$$

- It is not too hard to show that this produces the **same sequence** of matrices  $\mathbf{A}_k$  as the QR algorithm.

# Why the QR algorithm works

- Summary: The columns of  $\tilde{\mathbf{Q}}_k$  converge to the eigenvectors, and

$$\mathbf{A}_k = \tilde{\mathbf{Q}}_k^* \mathbf{A} \tilde{\mathbf{Q}}_k.$$

- We can recognize the above as a matrix of Rayleigh quotients, which for diagonalizable matrices

$$(\mathbf{A}_k)_{ij} = \tilde{\mathbf{q}}_i^* \mathbf{A} \tilde{\mathbf{q}}_j \rightarrow \lambda_i \delta_{ij} = \begin{cases} \lambda_i & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$

showing that (under suitable assumptions):

$$\mathbf{A}_k \rightarrow \mathbf{\Lambda}$$

- It can also be shown that

$$\tilde{\mathbf{Q}}_k = \mathbf{Q}_1 \mathbf{Q}_2 \cdots \mathbf{Q}_k \rightarrow \mathbf{X}$$

# More on $QR$ algorithm

- The convergence of the basic  $QR$  algorithm is closely related to that of the power method: It is only fast if all eigenvalues are **well-separated**.
- For more general (non-diagonalizable) matrices in complex arithmetic, the algorithm converges to the **Schur decomposition**  $\mathbf{A} = \mathbf{U}\mathbf{T}\mathbf{U}^*$ ,

$$\mathbf{A}_k \rightarrow \mathbf{T} \text{ and } \tilde{\mathbf{Q}}_k \rightarrow \mathbf{U}.$$

- It is possible to implement the algorithm entirely using real arithmetic (no complex numbers).
- There are several key improvements to the basic method that make this work in practice: **Hessenberg matrices** for faster  $QR$  factorization, **shifts** and **deflation** for acceleration.
- There are other sophisticated algorithms as well, such as the **divide-and-conquer algorithm**, and the best are implemented in the library LAPACK (MATLAB).