

Eigenvalues and Singular Values Power Method and MATLAB use

Aleksandar Donev
Courant Institute, NYU¹
donev@courant.nyu.edu

¹Course MATH-UA.0252/MA-UY_4424, Spring 2021

Spring 2021

Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Principal Component Analysis (PCA)

Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Principal Component Analysis (PCA)

The need for iterative algorithms

- The eigenvalues are roots of the **characteristic polynomial** of **A**, which is generally of order n .
- According to Abel's theorem, there is no closed-form (rational) solution for $n \geq 5$.

All eigenvalue algorithms must be iterative!

- There is an important distinction between iterative methods to:
 - Compute **all eigenvalues** (similarity transformations). These are based on dense-matrix factorizations such as the QR factorization, with total cost $O(n^3)$.
 - Compute **only one or a few eigenvalues**, typically the smallest or the largest one (e.g., power method). These are similar to iterative methods for solving linear systems.

Sparse Matrices

- Recall that for a diagonalizable matrix

$$\mathbf{A}^n = \mathbf{X}\mathbf{\Lambda}^n\mathbf{X}^{-1}$$

and **assume well-separated eigenvalues** $|\lambda_1| > |\lambda_2| \geq |\lambda_3| \cdots |\lambda_n|$, and that the columns of \mathbf{X} are normalized, $\|\mathbf{x}_j\| = 1$.

- For sparse matrices we sometimes only need to know a **few of the eigenvalues**/vectors, not all of them.
- Notably, knowing the eigenvector corresponding to the **smallest and largest (in magnitude) eigenvalues** is often most important (see Google Page Rank algorithm).

Iterative Method

- Any **initial guess** vector \mathbf{q}_0 can be represented in the linear basis formed by the eigenvectors

$$\mathbf{q}_0 = \mathbf{X}\mathbf{a}$$

- Recall iterative methods for linear systems: **Multiply a vector with the matrix \mathbf{A}** many times:

$$\mathbf{q}_{k+1} = \mathbf{A}\mathbf{q}_k$$

$$\mathbf{q}_n = \mathbf{A}^n \mathbf{q}_0 = (\mathbf{X}\mathbf{\Lambda}^n \mathbf{X}^{-1}) \mathbf{X}\mathbf{a} = \mathbf{X}(\mathbf{\Lambda}^n \mathbf{a})$$

Power Method

- As $n \rightarrow \infty$, the **eigenvalue of largest modulus** λ_0 will dominate,

$$\mathbf{\Lambda}^n = \lambda_1^n \text{Diag} \left\{ 1, \left(\frac{\lambda_2}{\lambda_1} \right)^n, \dots \right\} \rightarrow \text{Diag} \{ \lambda_1^n, 0, \dots, 0 \}$$

$$\mathbf{q}_n = \mathbf{X} (\mathbf{\Lambda}^n \mathbf{a}) \rightarrow \lambda_1^n \mathbf{X} \begin{bmatrix} a_1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \lambda_1^n \mathbf{x}_1$$

- Therefore the **normalized iterates** converge to the eigenvector:

$$\tilde{\mathbf{q}}_n = \frac{\mathbf{q}_n}{\|\mathbf{q}_n\|} \rightarrow \mathbf{x}_1$$

- The **Rayleigh quotient** converges to the eigenvalue:

$$r_A(\mathbf{q}_n) = \frac{\mathbf{q}_n^* \mathbf{A} \mathbf{q}_n}{\mathbf{q}_n \cdot \mathbf{q}_n} = \tilde{\mathbf{q}}_n^* \mathbf{A} \tilde{\mathbf{q}}_n \rightarrow \lambda_1$$

Power Iteration

Start with an initial guess \mathbf{q}_0 , and then iterate:

- 1 Compute **matrix-vector product** and normalize it:

$$\mathbf{q}_k = \frac{\mathbf{A}\mathbf{q}_{k-1}}{\|\mathbf{A}\mathbf{q}_{k-1}\|}$$

- 2 Use Raleigh quotient to obtain **eigenvalue estimate**:

$$\hat{\lambda}_k = \mathbf{q}_k^* \mathbf{A} \mathbf{q}_k$$

- 3 **Test for convergence**: Evaluate the residual

$$\mathbf{r}_k = \mathbf{A}\mathbf{q}_k - \hat{\lambda}_k \mathbf{q}_k$$

and terminate if the residual norm is smaller than some tolerance, e.g., for tolerance $\epsilon \ll 1$,

$$\|\mathbf{r}_k\| \approx |\lambda_1 - \hat{\lambda}_k| \leq \epsilon \hat{\lambda}_k.$$

Eigenvalues in MATLAB

- The **Schur decomposition** is provided by $[U, T] = \text{schur}(A)$.
- In MATLAB, sophisticated variants of the **QR algorithm** (LAPACK library) are implemented in the function *eig*:

$$\Lambda = \text{eig}(A)$$

$$[X, \Lambda] = \text{eig}(A)$$

- For large or sparse matrices, iterative methods based on the **Arnoldi iteration** (ARPACK library), can be used to obtain a few of the largest eigenvalues:

$$\Lambda = \text{eigs}(A, n_{\text{eigs}})$$

$$[X, \Lambda] = \text{eigs}(A, n_{\text{eigs}})$$

Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Principal Component Analysis (PCA)

Sensitivity (conditioning) of the SVD

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^*$$

- Since unitary matrices have unit 2-norm,

$$\|\delta\mathbf{\Sigma}\|_2 \approx \|\delta\mathbf{A}\|_2.$$

- The SVD computation is always **perfectly well-conditioned**!
- However, this refers to absolute errors: The **relative error** of small singular values will be large.
- The **power of the SVD** lies in the fact that it always exists and can be computed stably...but it is somewhat **expensive to compute**.

Computing the SVD

- The SVD can be computed by performing an eigenvalue computation for the **normal matrix** $\mathbf{A}^* \mathbf{A}$ (a positive-semidefinite matrix).
- This squares the condition number for small singular values and is **not numerically-stable**.
- Instead, modern algorithms use an algorithm based on computing eigenvalues / eigenvectors using the QR factorization.
- The cost of the calculation is $\sim O(mn^2)$, of the same order as eigenvalue calculation if $m \sim n$.

Reduced SVD

The **full (standard) SVD**

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^* = \sum_{i=1}^p \sigma_i \mathbf{u}_i \mathbf{v}_i^*$$

$$[m \times n] = [m \times m] [m \times n] [n \times n],$$

is in practice often computed in **reduced (economy) SVD** form, where $\mathbf{\Sigma}$ is $[p \times p]$:

$$[m \times n] = [m \times n] [n \times n] [n \times n] \quad \text{for } m > n$$

$$[m \times n] = [m \times m] [m \times m] [m \times n] \quad \text{for } n > m$$

This contains all the information as the full SVD but can be **cheaper to compute** if $m \gg n$ or $m \ll n$.

In MATLAB

- $[U, \Sigma, V] = \text{svd}(A)$ for **full SVD**, computed using a QR-like method.
- $[U, \Sigma, V] = \text{svd}(A, 'econ')$ for **economy SVD**.
- The **least-squares solution** for square, overdetermined, underdetermined, or even rank-deficient systems can be computed using *svd* or *pinv* (pseudo-inverse, see homework).
- The q largest singular values and corresponding approximation can be computed efficiently for **sparse matrices** using

$$[U, \Sigma, V] = \text{svds}(A, q).$$

Outline

- 1 Eigenvalue Problems
- 2 Singular Value Decomposition
- 3 Principal Component Analysis (PCA)

Low-rank approximations

- The SVD is a decomposition into **rank-1 outer product matrices**:

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^* = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^* = \sum_{i=1}^r \mathbf{A}_i$$

- The rank-1 components \mathbf{A}_i are called **principal components**, the most important ones corresponding to the larger σ_i .
- Ignoring all singular values/vectors except the first q , we get a **low-rank approximation**:

$$\mathbf{A} \approx \hat{\mathbf{A}}_q = \mathbf{U}_q \mathbf{\Sigma}_q \mathbf{V}_q^* = \sum_{i=1}^q \sigma_i \mathbf{u}_i \mathbf{v}_i^*.$$

- Theorem: This is the **best approximation** of rank- q in the Euclidian and Frobenius norm:

$$\left\| \mathbf{A} - \hat{\mathbf{A}}_q \right\|_2 = \sigma_{q+1}$$

Applications of SVD/PCA

- **Statistical analysis** (e.g., DNA microarray analysis, clustering).
- Data **compression** (e.g., image compression, explained next).
- **Feature extraction**, e.g., face or character recognition (see Eigenfaces on Wikipedia).
- **Latent semantic indexing** for context-sensitive searching (see Wikipedia).
- **Noise reduction** (e.g., weather prediction).
- One example concerning language analysis given in homework.

Image Compression

```
>> A=rgb2gray(imread('basket.jpg'));  
>> imshow(A);  
>> [U,S,V]=svd(double(A));  
>> r=25; % Rank-r approximation  
>> Acomp=U(:,1:r)*S(1:r,1:r)*(V(:,1:r))';  
>> imshow(uint8(Acomp));
```

Compressing an image of a basket

We used only 25 out of the ~ 400 singular values to construct a rank 25 approximation:





Principal Component Analysis

- **Principal Component Analysis (PCA)** is a term used for low-rank approximations in statistical analysis of data.
- Consider having m empirical data points or **observations** (e.g., daily reports) of n **variables** (e.g., stock prices), and put them in a **data matrix** $\mathbf{A} = [m \times n]$.
- Assume that each of the variables has **zero mean**, that is, the empirical mean has been subtracted out.
- It is also useful to choose the units of each variable (normalization) so that the **variance is unity**.
- We would like to find an **orthogonal transformation** of the original variables that accounts for as much of the variability of the data as possible.
- Specifically, the first principal component is the direction along which the variance of the data is largest.

PCA and Variance

PCA applied to an ellipsoidally shaped point cloud



more information: www.joyofdata.de/blog/illustration-of-principal-component-analysis-pca

PCA and SVD

- The **covariance matrix** of the data tells how correlated different pairs of variables are:

$$\mathbf{C} = \mathbf{A}^T \mathbf{A} = [n \times n]$$

- The largest eigenvalue of \mathbf{C} is the direction (line) that minimizes the sum of squares of the distances from the points to the line, or equivalently, **maximizes the variance** of the data projected onto that line.
- The SVD of the data matrix is $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^*$.
- The eigenvectors of \mathbf{C} are in fact the columns of \mathbf{V} , and the eigenvalues of \mathbf{C} are the squares of the singular values,

$$\mathbf{C} = \mathbf{A}^T \mathbf{A} = \mathbf{V}\mathbf{\Sigma} (\mathbf{U}^* \mathbf{U}) \mathbf{\Sigma} \mathbf{V}^* = \mathbf{V}\mathbf{\Sigma}^2 \mathbf{V}^*.$$

Note: the eigenvalues values are necessarily real and positive since \mathbf{C} is positive semi-definite.

Dimensionality reduction via PCA

